

▼ DeepLense: Exploring Transformers

```
from google.colab import drive
drive.mount('/content/gdrive')
```

```
Mounted at /content/gdrive
```

```
!tar xzf gdrive/MyDrive/lenses.tar.gz
```

```
!pip install pytorch_lightning
```

```
Requirement already satisfied: pytorch_lightning in /usr/local/lib/python3.7
Requirement already satisfied: torchmetrics>=0.4.1 in /usr/local/lib/python3
Requirement already satisfied: PyYAML>=5.4 in /usr/local/lib/python3.7/dist-
Requirement already satisfied: typing-extensions>=4.0.0 in /usr/local/lib/py
Requirement already satisfied: pyDeprecate<0.4.0,>=0.3.1 in /usr/local/lib/p
Requirement already satisfied: numpy>=1.17.2 in /usr/local/lib/python3.7/dis
Requirement already satisfied: torch>=1.8.* in /usr/local/lib/python3.7/dist
Requirement already satisfied: packaging>=17.0 in /usr/local/lib/python3.7/d
Requirement already satisfied: tensorboard>=2.2.0 in /usr/local/lib/python3.
Requirement already satisfied: tqdm>=4.41.0 in /usr/local/lib/python3.7/dist
Requirement already satisfied: fsspec[http]!=2021.06.0,>=2021.05.0 in /usr/l
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-pac
Requirement already satisfied: aiohttp in /usr/local/lib/python3.7/dist-pack
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /usr/local/lib/py
Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /usr
Requirement already satisfied: protobuf>=3.6.0 in /usr/local/lib/python3.7/d
Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /usr/loca
Requirement already satisfied: absl-py>=0.4 in /usr/local/lib/python3.7/dist
Requirement already satisfied: grpcio>=1.24.3 in /usr/local/lib/python3.7/di
Requirement already satisfied: setuptools>=41.0.0 in /usr/local/lib/python3.
Requirement already satisfied: werkzeug>=0.11.15 in /usr/local/lib/python3.7
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/pytho
Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /usr/local/l
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.7/d
Requirement already satisfied: wheel>=0.26 in /usr/local/lib/python3.7/dist-
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.7/dis
Requirement already satisfied: cachetools<5.0,>=2.0.0 in /usr/local/lib/pyth
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/pytho
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/py
Requirement already satisfied: importlib-metadata>=4.4 in /usr/local/lib/pyt
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-pa
Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /usr/local/lib/python
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /u
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.7/d
Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.7/di
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.7/dis
Requirement already satisfied: async-timeout<5.0,>=4.0.0a3 in /usr/local/lib
Requirement alreadyv satisfied: frozenlist>=1.1.1 in /usr/local/lib/pvthon3.7
```

✓ 1m 7s completed at 7:24 PM

● ✕

Requirement already satisfied: charset-normalizer<3.0,>=2.0 in /usr/local/li
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3
Requirement already satisfied: asyncctest==0.13.0 in /usr/local/lib/python3.7

Standard libraries

import os

import shutil

import numpy as np

import random

import math

import json

from functools import partial

from PIL import Image

Imports for plotting

import matplotlib.pyplot as plt

%matplotlib inline

PyTorch

import torch

import torch.nn as nn

import torch.nn.functional as F

import torch.utils.data as data

import torch.optim as optim

Torchvision

import torchvision

from torchvision import transforms

from torchvision.utils import make_grid

from torchvision.datasets import DatasetFolder, ImageFolder

from torch.utils.data import DataLoader

Imports for ROC AUC

from sklearn.preprocessing import LabelBinarizer

from sklearn.metrics import roc_curve, roc_auc_score, auc

from itertools import cycle

Imports for PyTorch Lightning

import pytorch_lightning as pl

from pytorch_lightning.callbacks import LearningRateMonitor, ModelCheckpoint

pl.seed_everything(42)

device = torch.device("cuda:0") if torch.cuda.is_available() else torch.device("c
print("Device:", device)

Global seed set to 42

Device: cuda:0

Kaggle kernel

LENSES_DATASET_PATH = "/content/lenses" # Should point to the root of the dataset

BATCH_SIZE = 16

TEST_SPLIT=0.1

```

VAL_SPLIT=0.01
INPUT_HEIGHT = 150
INPUT_WIDTH = 150
TRAIN='train'
TEST='test'
VAL='val'

lenses_files = []
for folder in {'sub', 'no_sub'}:
    for file in os.listdir(os.path.join(LENSES_DATASET_PATH, folder)):
        if file.endswith(".jpg"):
            lenses_files.append(os.path.join(LENSES_DATASET_PATH, folder, file))

def copy_images(imagePaths, folder):
    if not os.path.exists(folder):
        os.makedirs(folder)
    for path in imagePaths:
        imageName = path.split(os.path.sep)[-1]
        label = path.split(os.path.sep)[-2]
        labelFolder = os.path.join(folder, label)
        if not os.path.exists(labelFolder):
            os.makedirs(labelFolder)
        destination = os.path.join(labelFolder, imageName)
        shutil.copy(path, destination)

np.random.shuffle(lenses_files)
valPathsLen = int(len(lenses_files) * VAL_SPLIT)
testPathsLen = int(len(lenses_files) * TEST_SPLIT)
trainPathsLen = len(lenses_files) - valPathsLen - testPathsLen
print(f"Train : {trainPathsLen}, Test: {testPathsLen}, Val:{valPathsLen}")

    Train : 8900, Test: 1000, Val:100

testInd = trainPathsLen + testPathsLen

trainPaths = lenses_files[:trainPathsLen]
testPaths = lenses_files[trainPathsLen:testInd]
valPaths = lenses_files[testInd:]

copy_images(trainPaths, TRAIN)
copy_images(testPaths, TEST)
copy_images(valPaths, VAL)

trainTransforms = transforms.Compose([
    transforms.Resize(size=(224, 224)),
    transforms.ToTensor()
])
testTransforms = transforms.Compose([
    transforms.Resize(size=(224, 224)),
    transforms.ToTensor()
])

```

```

transforms.Compose([
    transforms.ToTensor()
])

trainDataset = ImageFolder(root=TRAIN, transform=trainTransforms)
testDataset = ImageFolder(root=TEST, transform=testTransforms)
valDataset = ImageFolder(root=VAL, transform=testTransforms)

print(f"[INFO] Training dataset contains {len(trainDataset)} samples.")
print(f"[INFO] Test dataset contains {len(testDataset)} samples.")
print(f"[INFO] Validation dataset contains {len(valDataset)} samples.")

[INFO] Training dataset contains 8900 samples.
[INFO] Test dataset contains 1000 samples.
[INFO] Validation dataset contains 100 samples.

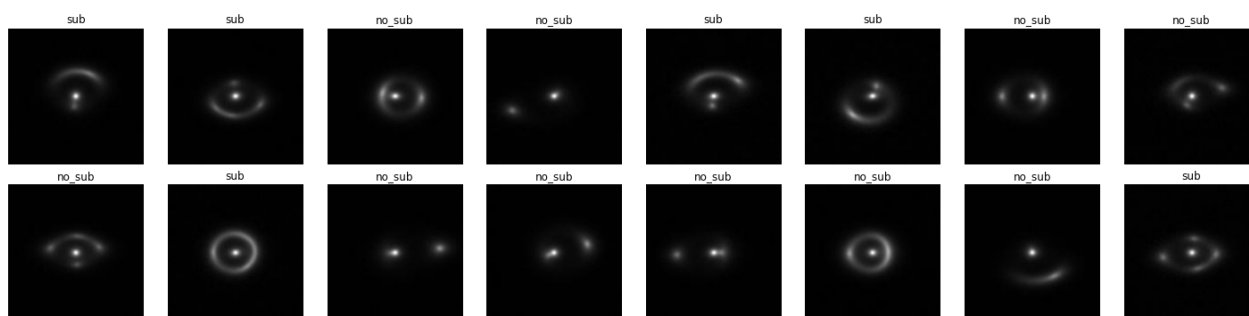
trainDataLoader = DataLoader(trainDataset,
                              batch_size=BATCH_SIZE, shuffle=True, num_workers=2)
testDataLoader = DataLoader(testDataset, shuffle=False, drop_last=False, batch_size=BATCH_SIZE)
valDataLoader = DataLoader(valDataset, shuffle=False, drop_last=False, batch_size=BATCH_SIZE)

def visualize_batch(batch, classes, dataset_type):
    fig = plt.figure("{} batch".format(dataset_type), figsize=(20, 5))
    for i in range(0, BATCH_SIZE):
        ax = plt.subplot(2, 8, i+1)
        image = batch[0][i].cpu().numpy()
        image = image.transpose((1, 2, 0))
        image = (image * 255.0).astype("uint8")
        idx = batch[1][i]
        label = classes[idx]
        plt.imshow(image)
        plt.title(label)
        plt.axis("off")
    plt.tight_layout()
    plt.show()

trainBatch = next(iter(trainDataLoader))
print(trainBatch[0].shape)
visualize_batch(trainBatch, trainDataset.classes, "train")

```

```
torch.Size([16, 3, 224, 224])
```

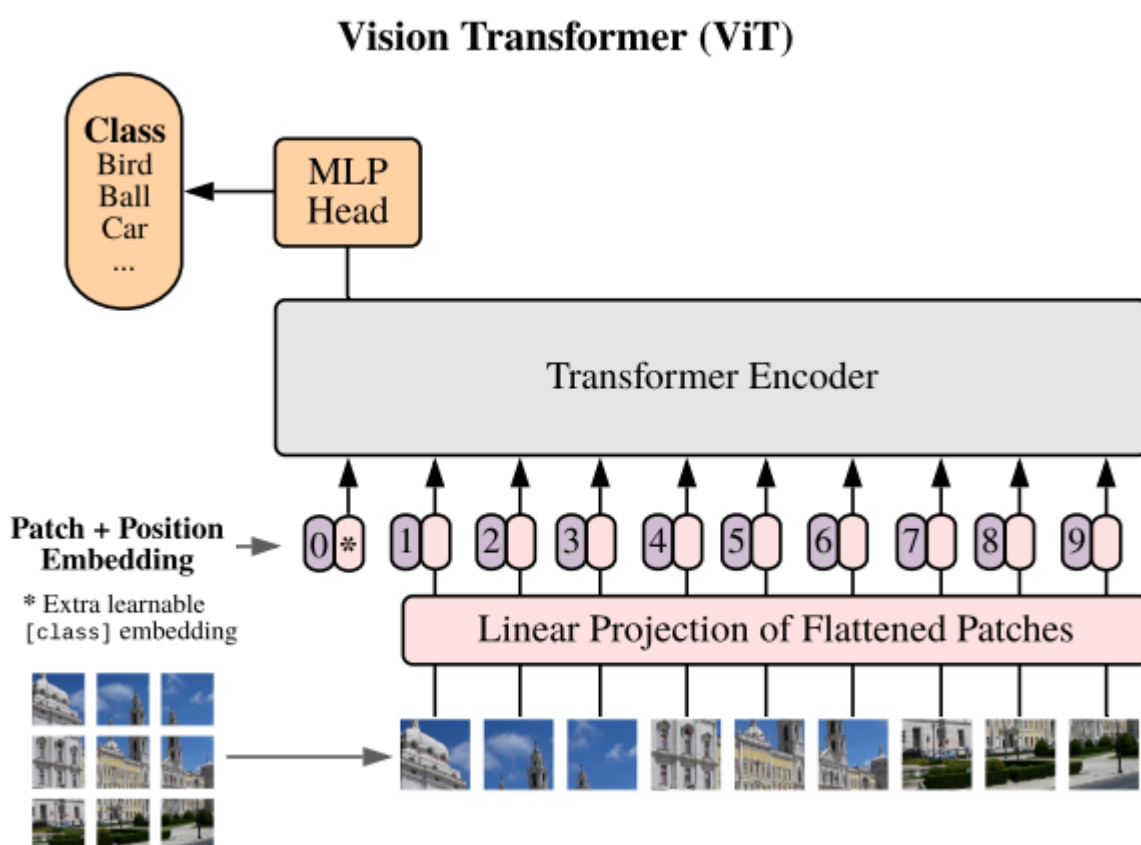


Breaking image into multiple patches and the flattening them.m

```
def img_to_patch(x, patch_size, flatten_channels=True):
    B, C, H, W = x.shape
    #x = x.reshape(B, C, H//patch_size, patch_size, W//patch_size, patch_size)
    x = x.reshape(B, C, torch.div(H, patch_size, rounding_mode='trunc'), patch_s:

    x = x.permute(0, 2, 4, 1, 3, 5)
    x = x.flatten(1,2)
    if flatten_channels:
        x = x.flatten(2,4)
    return x
```

Vision Transformer



Patch Embeddings

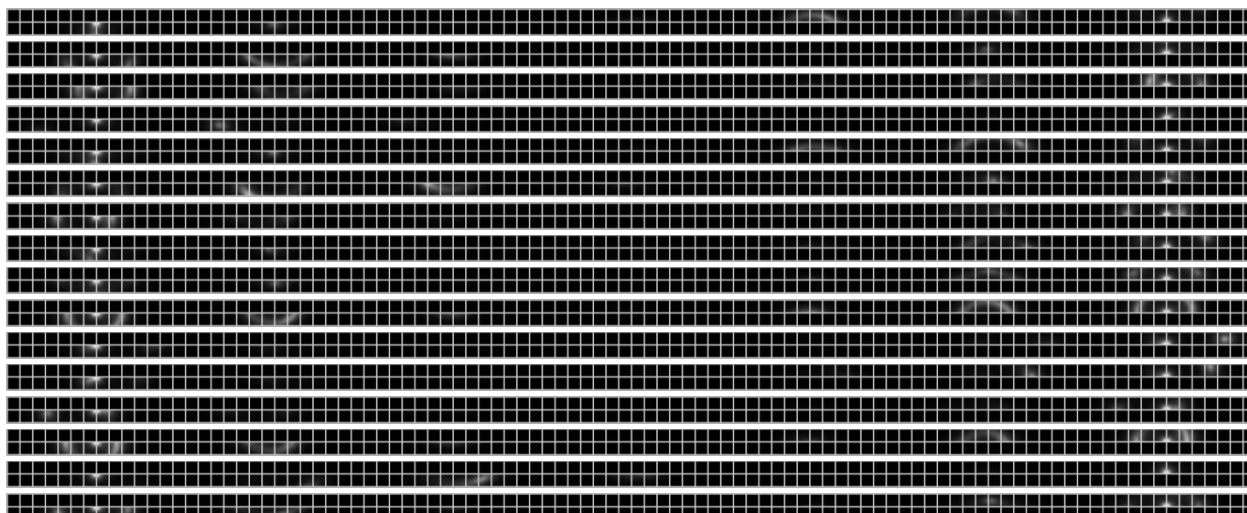
```
img_patches = img_to_patch(trainBatch[0], patch_size=16, flatten_channels=False)

fig, ax = plt.subplots(trainBatch[0].shape[0], 1, figsize=(40,10))

print('Display patch embedding result: ')
# 224x224 images, split as 16x16 patches so 196 patches in total, displaying 98

for i in range(trainBatch[0].shape[0]):
    img_grid = make_grid(img_patches[i], nrow=98, normalize=True, pad_value=0.8)
    img_grid = img_grid.permute(1, 2, 0)
    ax[i].imshow(img_grid)
    ax[i].axis('off')
plt.show()
plt.close()
```

Display patch embedding result:



The Vision Transformer model

```

class AttentionBlock(nn.Module):
    """
    embed_dim - dims of the input and attn of feature vectors
    hidden_dim - dims of hidden layer in FFN i.e 2 x embed_dim
    num_heads - number of heads in MHA block
    dropout - amount of dropout to apply in FFN
    """
    def __init__(self, embed_dim, hidden_dim, num_heads, dropout=0.0):
        super().__init__()

        self.layer_norm_1 = nn.LayerNorm(embed_dim)
        self.attn = nn.MultiheadAttention(embed_dim, num_heads)
        self.layer_norm_2 = nn.LayerNorm(embed_dim)
        self.linear = nn.Sequential(
            nn.Linear(embed_dim, hidden_dim),
            nn.GELU(),
            nn.Dropout(dropout),
            nn.Linear(hidden_dim, embed_dim),
            nn.Dropout(dropout)
        )

    def forward(self, x):
        inp_x = self.layer_norm_1(x)
        x = x + self.attn(inp_x, inp_x, inp_x)[0]
        x = x + self.linear(self.layer_norm_2(x))
        return x


class VisionTransformer(nn.Module):
    """
    embed_dim - dims of the input feature vectors
    hidden_dim - dims of the hidden layer in the FFN within transformer
    num_channels - num of channels of the input i.e. 3 in our case of RGB im
    num_heads - num of heads to use in the MHA block
    num_layers - num of layers in transformer
    num_classes - num of classes to predict
    patch_size - num of pixels that the patches have per dim
    num_patches - max number of patches of an image
    dropout - amount of dropout to apply in the FFN and on the input encoding
    """
    def __init__(self, embed_dim, hidden_dim, num_channels, num_heads, num_layer:
        super().__init__()

        self.patch_size = patch_size

        self.input_layer = nn.Linear(num_channels*(patch_size**2), embed_dim)
        self.transformer = nn.Sequential(*[AttentionBlock(embed_dim, hidden_dim,
        self.mlp_head = nn.Sequential(
            nn.LayerNorm(embed_dim),
            nn.Linear(embed_dim, num_classes)
        )
        self.dropout = nn.Dropout(dropout)

```

```

self.cls_token = nn.Parameter(torch.randn(1, 1, embed_dim))
self.pos_embedding = nn.Parameter(torch.randn(1, 1+num_patches, embed_dim))

```

```

def forward(self, x):
    x = img_to_patch(x, self.patch_size)
    B, T, _ = x.shape
    x = self.input_layer(x)

    cls_token = self.cls_token.repeat(B, 1, 1)
    x = torch.cat([cls_token, x], dim=1)
    x = x + self.pos_embedding[:, :T+1]

    x = self.dropout(x)
    x = x.transpose(0, 1)
    x = self.transformer(x)

    cls = x[0]
    out = self.mlp_head(cls)
    return out

```

We will be using PyTorch's Lightning module to organize our model code.

```

class ViT(pl.LightningModule):
    def __init__(self, model_kwargs, lr):
        super().__init__()
        self.save_hyperparameters()
        self.model = VisionTransformer(**model_kwargs)
        self.example_input_array = next(iter(trainDataLoader))[0]
        self.predictions = [] # Actual predictions for ROC AUC

    def forward(self, x):
        return self.model(x)

    def configure_optimizers(self):
        optimizer = optim.AdamW(self.parameters(), lr=self.hparams.lr)
        lr_scheduler = optim.lr_scheduler.MultiStepLR(optimizer, milestones=[100],
        return [optimizer], [lr_scheduler])

    def _calculate_loss(self, batch, mode="train"):
        imgs, labels = batch
        preds = self.model(imgs)
        loss = F.cross_entropy(preds, labels)
        acc = (preds.argmax(dim=-1) == labels).float().mean()

        if mode == 'test':
            self.predictions.append(preds.argmax(dim=-1))

        self.log(f'{mode}_loss', loss, )
        self.log(f'{mode}_acc', acc)
        return loss

    def training_step(self, batch, batch_idx):

```



```

def training_step(self, batch, batch_idx):
    loss = self._calculate_loss(batch, mode="train")
    return loss

def validation_step(self, batch, batch_idx):
    self._calculate_loss(batch, mode="val")

def test_step(self, batch, batch_idx):
    self._calculate_loss(batch, mode="test")

CHECKPOINT_PATH = "./saved_models/"

def train_model(**kwargs):
    trainer = pl.Trainer(default_root_dir=os.path.join(CHECKPOINT_PATH, "ViT"),
                        gpus=1 if str(device)=="cuda:0" else 0,
                        max_epochs=180,
                        callbacks=[ModelCheckpoint(save_weights_only=True, mode=
                                LearningRateMonitor("epoch"))],
                        enable_progress_bar=False,
                        log_every_n_steps=4)

    #trainer.logger._log_graph = True # If True, we plot the computation
    #trainer.logger._default_hp_metric = None # Optional logging argument that we

    pl.seed_everything(42, workers=True) # To be reproducible
    model = ViT(**kwargs)
    trainer.fit(model=model, train_dataloaders=trainDataLoader, val_dataloaders=valDataLoader)

    # Using PyTorch Lightning we can load a revised model and use it as the best
    val_result = trainer.validate(model, dataloaders=valDataLoader, verbose=False)
    test_result = trainer.test(model, dataloaders=testDataLoader, verbose=False)

    result = {"test": test_result[0]["test_acc"], "val": val_result[0]["val_acc"]}

    return model, result

EMBED_DIM=256
HIDDEN_DIM=512
RESIZE_IMG=224
NUM_HEADS=8
NUM_LAYERS=6
PATCH_SIZE=16
NUM_CHANNELS=3
NUM_PATCHES=int((RESIZE_IMG*RESIZE_IMG) / (PATCH_SIZE*PATCH_SIZE))
NUM_CLASSES=2
DROPOUT=0.2

```

Training the model

```
model, results = train_model(model_kwargs={
```

```

        'embed_dim': EMBED_DIM,
        'hidden_dim': HIDDEN_DIM,
        'num_heads': NUM_HEADS,
        'num_layers': NUM_LAYERS,
        'patch_size': PATCH_SIZE,
        'num_channels': NUM_CHANNELS,
        'num_patches': NUM_PATCHES,
        'num_classes': NUM_CLASSES,
        'dropout': DROPOUT
    }, lr=3e-4)

print("ViT results", results)

```

```

GPU available: True, used: True
TPU available: False, using: 0 TPU cores
IPU available: False, using: 0 IPUs
HPU available: False, using: 0 HPUs
Global seed set to 42
Missing logger folder: saved_models/ViT/lightning_logs
LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]

```

	Name	Type	Params	In sizes	Out sizes
0	model	VisionTransformer	3.4 M	[16, 3, 224, 224]	[16, 2]
3.4 M		Trainable params			
0		Non-trainable params			
3.4 M		Total params			
13.645		Total estimated model params size (MB)			
LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]					
LOCAL_RANK: 0 - CUDA_VISIBLE_DEVICES: [0]					
ViT results {'test': 0.968999981880188, 'val': 0.9800000190734863}					

ROC AUC Score

```

class LB(LabelBinarizer):
    def transform(self, y):
        Y = super().transform(y)
        if self.y_type_ == 'binary':
            return np.hstack((Y, 1-Y))
        else:
            return Y

    def inverse_transform(self, Y, threshold=None):
        if self.y_type_ == 'binary':
            return super().inverse_transform(Y[:, 0], threshold)
        else:
            return super().inverse_transform(Y, threshold)

y_score = torch.cat(model.predictions).cpu().detach().numpy()

y_test = []
for _, labels in testDataLoader:

```

```

y_test.append(labels.cpu().detach().numpy())

y_test = np.concatenate(y_test)

lb = LB()
y_test = lb.fit_transform(y_test)
y_score = lb.fit_transform(y_score)

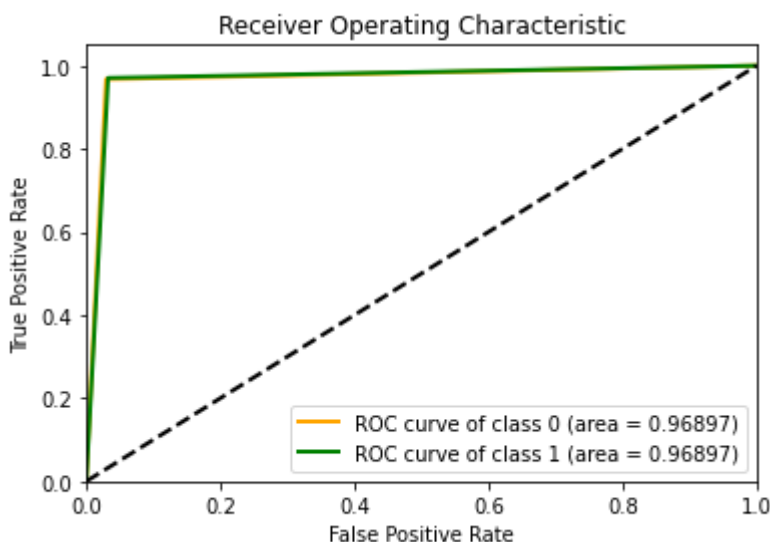
n_classes = y_test.shape[1]
fpr = dict()
tpr = dict()
roc_auc = dict()

for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])
colors = ['orange', 'green']

for i, color in zip(range(n_classes), colors):
    plt.plot(fpr[i], tpr[i], color=color, lw=2, label='ROC curve of class {0} (a:

plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()

```



Save Model

```
torch.save(model.state_dict(), 'st5_model.pth')
```

